

Distributed PageRank

Aaron Myers, Megan Ruthven

October 30, 2014

1 Abstract

1.1 Problem Description

We will focus on developing and testing iterative methods for distributed page rank. Traditionally, PageRank has been computed as the principal eigenvector of a Markov chain probability transition matrix using a simple power iteration algorithm. We can think of computing PageRank as solving a system of linear equations, and apply iterative solvers to obtain the solution. Parallelization is an important factor in choosing a particular iterative linear solver. We will attempt to develop a distributed algorithm for computing PageRank (both global and personalized) of very large web graphs using MPI. Load balancing is important factor in designing our distributed algorithm. We will consider many strategies, for example, distributing load based on the topology of the graph.

1.2 Proposed Initial Strategies

1. Krylov Subspace Methods

In the typical power iteration, we compute Pb, P^2b, P^3b, \dots where b is a random vector, ultimately converging on the principle eigenvector. However, much computation is wasted by using only the final result $A^{n-1}b$. We will attempt to use the Krylov Matrix:

$$K_n = [b \quad Ab \quad A^2b \quad A^3b \dots A^{n-1}b] \quad (1)$$

The columns are not orthogonal, but we can extract an orthogonal basis and use this to find an approximation to the eigenvectors and understand how 'far apart' the eigenvectors are so we can estimate the iterations needed to converge. Other Krylov methods include GMRES, Jacobi, Block Jacobi, etc.

2. ADMM

Here we proposed applying ADMM as a way to parallelize some iterative method (possible the power iteration or a Krylov method) to converge onto a solution faster than simply parallelizing the matrix multiply. If ADMM proves unsuccessful, we may attempt to distribute the computing load based on some clustering or tree method.

3. Low Rank Approximation to P - Secondary

To simplify the matrix multiply, we will attempt some form of low rank approximation for P . There are many methods proposed here, Frequent Directions, Sampling, Random SVD, Hashing, Random projection, etc. If we could simplify the matrix vector multiply in the power iteration and combine with some form of ADMM, we may be able to significantly reduce computation time. This will likely not be our primary focus, but may help improve overall efficiency slightly.

4. Summary

Finally we would apply our proposed approach to a large graph problem and compare other methods as well as some naive approach.